

Introduction to Operating System

Apurba Sarkar

IEST Shibpur

July 21, 2016

Overview

- What is an operating system?
- Operating systems history
- Computer system and operating system structure

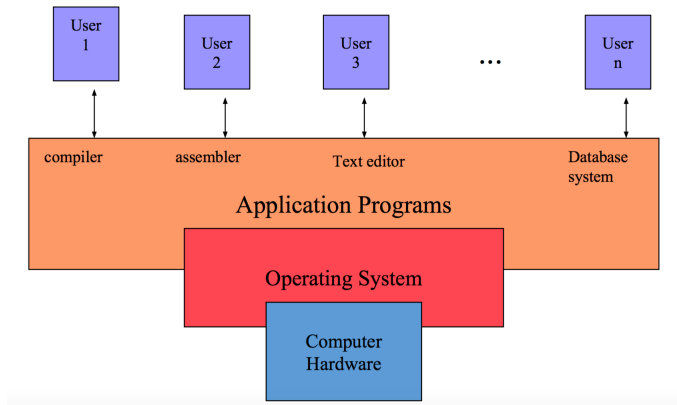
What is an operating system?

- OS is the software that acts an intermediary between the user applications and computer hardware.

Computer System Components

- Hardware
 - Provides basic computing resources (CPU, memory, I/O devices).
- Operating System
 - Controls and coordinates the use of hardware among application programs
- Application Programs
 - Solve computing problems of users (compilers, database systems, video games, business programs such as banking software).
- Users
 - People, machines, other computers

Abstract View of System



Operating system rolls

- Referee:
 - Resource allocation among users, applications
 - Isolation of different users, applications from each other
 - Communication between users, applications
- Illusionist
 - Files can grow (nearly) arbitrarily large
 - Files persist even when the machine crashes in the middle of a save
- Glue
 - Named directories, printf, . . .

Operating system challenges

- Reliability
 - Does the system do what it was designed to do?

Operating system challenges

- Reliability
 - Does the system do what it was designed to do?
- Availability
 - What portion of the time is the system working?

Operating system challenges

- Reliability
 - Does the system do what it was designed to do?
- Availability
 - What portion of the time is the system working?
- Security
 - Can the system be compromised by an attacker?

Operating system challenges

- Reliability
 - Does the system do what it was designed to do?
- Availability
 - What portion of the time is the system working?
- Security
 - Can the system be compromised by an attacker?
- Privacy
 - Whether data is accessible only to authorized users?

Operating system challenges contd.

- Performance
 - Latency/response time
 - How long does an operation take to complete?

Operating system challenges contd.

- Performance
 - Latency/response time
 - How long does an operation take to complete?
 - Throughput
 - How many operations can be done per unit of time?

Operating system challenges contd.

- Performance
 - Latency/response time
 - How long does an operation take to complete?
 - Throughput
 - How many operations can be done per unit of time?
 - Overhead
 - How much extra work is done by the OS?

Operating system challenges contd.

- Performance
 - Latency/response time
 - How long does an operation take to complete?
 - Throughput
 - How many operations can be done per unit of time?
 - Overhead
 - How much extra work is done by the OS?
 - Fairness
 - How equal is the performance received by different users?

Operating system challenges contd.

- Performance
 - Latency/response time
 - How long does an operation take to complete?
 - Throughput
 - How many operations can be done per unit of time?
 - Overhead
 - How much extra work is done by the OS?
 - Fairness
 - How equal is the performance received by different users?
 - Predictability
 - How consistent is the performance over time?

Why should I study Operating Systems?

- Need to understand interaction between the hardware and software
- Need to understand basic principles in the design of computer systems
 - efficient resource management, security, flexibility
- Above all of these, it enables you to do things that are difficult/impossible otherwise.


Operating systems history

- Early Operating Systems
- Simple Batch Systems
- Multiprogrammed Batch Systems
- Time-sharing Systems
- Personal and mobile Computer Systems

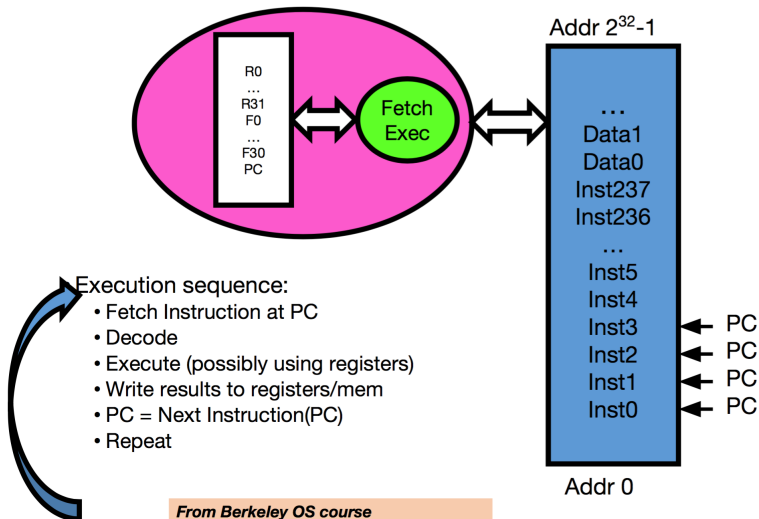
Computer System & OS Structures

- Computer System Organization
- Process abstraction and hardware protection
- System call and OS services
- Storage architecture
- OS organization
- OS Tasks
- Virtual Machines

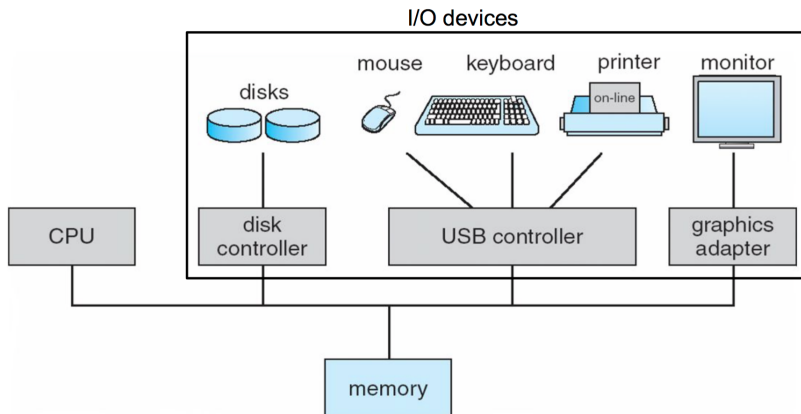
Computer System Organization



CPU Execution



Computer System Organization



I/O devices

- I/O devices and the CPU execute concurrently.
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer. I/O is from the device to local buffer of controller
- CPU moves data from/to main memory to/from the local buffers
- Device controller interrupts CPU on completion of I/O

Interrupts

- Interrupt transfers control to the interrupt service routine
 - Interrupt Service Routine: Segments of code that determine action to be taken for interrupt.
- Determining the type of interrupt
 - Polling: same interrupt handler called for all interrupts, which then polls all devices to figure out the reason for the interrupt
 - Interrupt Vector Table: different interrupt handlers will be executed for different interrupts.

Interrupt handling

- OS preserves the state of the CPU
 - stores registers and the program counter (address of interrupted instruction).
- Incoming interrupts are disabled while another interrupt is being processed to prevent a lost interrupt.

Different types of I/O processing for programs

- Synchronous I/O
 - After I/O is requested, wait until I/O is done. Program will be idle.
- Asynchronous I/O
 - After I/O is requested, control returns to user program without waiting for I/O completion.

Process Abstraction

- **Process**: an instance of a program, running with limited rights
 - **Thread**: a sequence of instructions within a process
 - Potentially many threads per process.
 - Address space: set of rights of a process
 - Memory that the process can access
 - Other permissions the process has (e.g., which system calls it can make, what files it can access)

How to limit process rights

Hardware Protection

- Dual Mode Operation
- Memory Protection
- CPU Protection
- I/O Protection

Should a process be able to execute any instructions


- No
 - Can alter system configuration
 - Can access unauthorized memory
 - Can access unauthorized I/O
 - etc.
- How to prevent?

Dual-mode operation

- Provide hardware support to differentiate between at least two modes of operation:
 - User mode – execution done on behalf of a user.
 - Kernel mode (monitor/supervisor/system mode) – execution done on behalf of operating system.
- “Privileged” instructions are only executable in the kernel mode
- Executing privileged instructions in the user mode “traps” into the kernel mode
 - Trap is a software generated interrupt caused either by an error or a user request.


Dual-mode operation contd.

- Mode bit added to computer hardware to indicate the current mode: kernel(0) or user(1).
- When an interrupt or trap occurs, hardware switches to kernel mode.



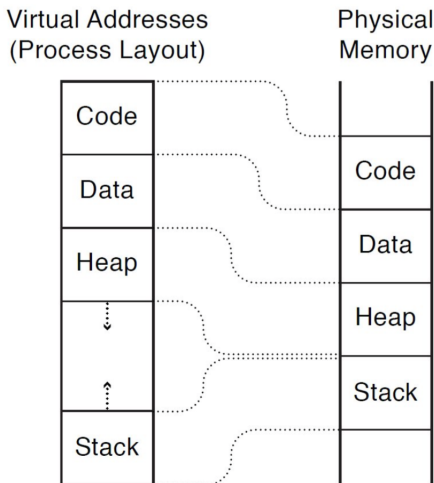
Process Address-Space

- Address space \Rightarrow the set of accessible addresses + state associated with them:
- For a 32-bit processor there are $2^{32} = 4$ billion addresses

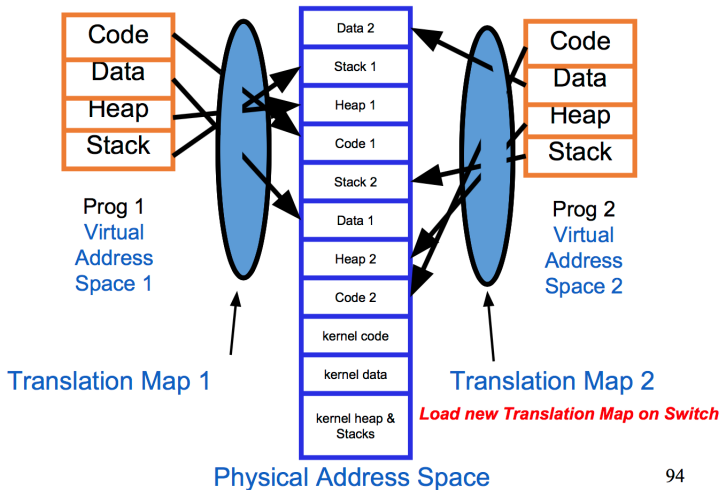


Virtual Address

- Provides the illusion of Separate Address Spaces

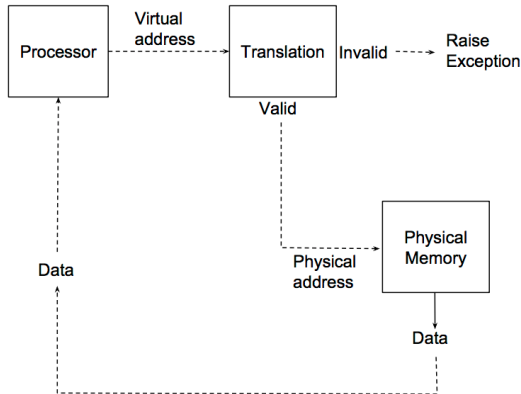


Virtual Address



94

Address translation and memory protection




Memory Protection

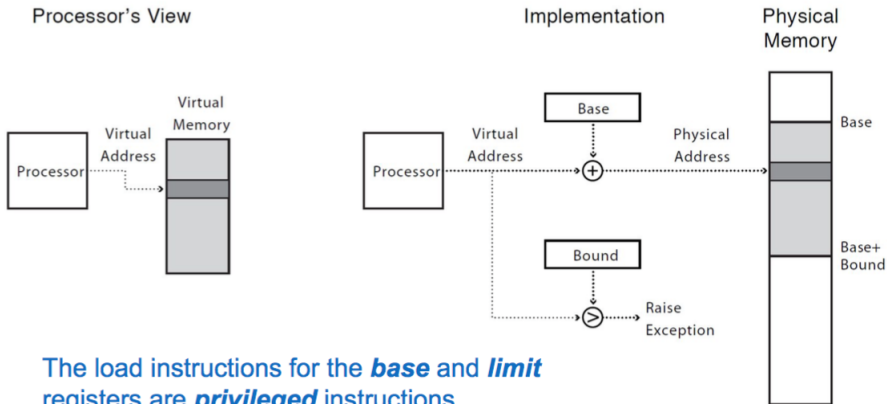
- Must provide memory protection at least for the interrupt vector and the interrupt service routines.
- When a process is running, only memory in that process address space must be accessible.
- When executing in kernel mode, the kernel has unrestricted access to all memory.

Memory Protection: base and bounds

- To provide memory protection, add two registers that determine the range of legal addresses a program may address.
 - Base Register - holds smallest legal physical memory address.
 - Limit register - contains the size of the range.
- Memory outside the defined range is protected.



Virtually Addressed Base and Bounds



CPU Protection

- How to prevent a process from executing indefinitely?

CPU Protection

- Timer - interrupts computer after specified period to ensure that OS maintains control.
 - Timer is decremented every clock tick.
 - When timer reaches a value of 0, an interrupt occurs.
- Timer is commonly used to implement time sharing.
- Timer is also used to compute the current time.
- Load timer is a privileged instruction.

I/O Protection


- All I/O instructions are privileged instructions.

A Question

- Given the I/O instructions are privileged, how do users perform I/O?
- Via system calls - the method used by a process to request action by the operating system.


System Calls

- User code can issue a syscall, which causes a trap
- Kernel handles the syscall.



System Calls

- Interface between running program and the OS.
 - Assembly language instructions (macros and subroutines)
 - Some higher level languages allow system calls to be made directly (e.g. C)
- Passing parameters between a running program and OS via registers, memory tables or stack.
- Linux has about 300 system calls
- `read()`, `write()`, `open()`, `close()`, `fork()`, `exec()`, `ioctl()`,

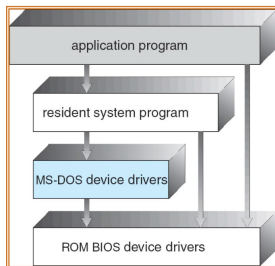


How are the Operating System organized?

- Simple
 - Only one or two levels of code
- Layered
 - Lower levels independent of upper levels
- Modular
 - Core kernel with Dynamically loadable modules
- Microkernel
 - OS built from many user-level processes

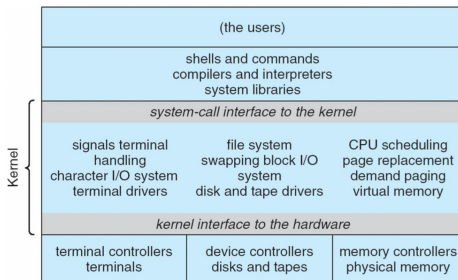
Os structure: Simple Approach

- MS-DOS - provides a lot of functionality in little space.
 - Not divided into modules, Interfaces and levels of functionality are not well separated.



Original UNIX System Structure


- Limited structuring, has 2 separable parts
 - Systems programs
 - shells and commands
 - compilers and interpreters
 - system libraries
 - Kernel
 - everything below system call interface and above physical hardware.
 - Filesystem, CPU scheduling, memory management.



Layered OS Structure

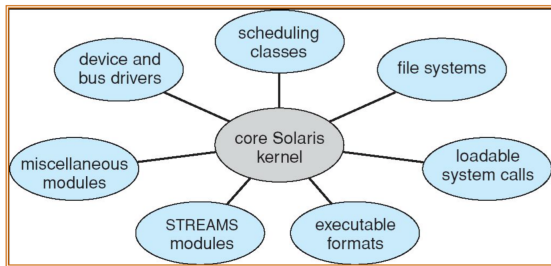
- OS divided into number of layers - bottom layer is hardware, highest layer is the user interface.
- Each layer uses functions and services of only lower- level layers.
- THE Operating System and Linux Kernel has successive layers of abstraction

Layered OS Structure



Modules-based Structure

- Most modern operating systems implement modules
 - Each core component is separate
 - Each talks to the others over known interfaces
 - Each is loadable as needed within the kernel
- Overall, similar to layers but with more flexibility
 - Linux makes extensive use of modules



Monolithic vs. Microkernel OS

- Monolithic OSES have large kernels with a lot of components
 - Linux, Windows, Mac
- Microkernels moves as much from the kernel into “user” space
 - Small core OS running at kernel level
 - OS Services built from many independent user-level processes
- Communication between modules with message passing
- Benefits
 - Easier to extend a microkernel
 - Easier to port OS to new architectures
 - More reliable (less code is running in kernel mode)
 - Fault Isolation (parts of kernel protected from other parts)
 - More secure

OS Task: Process Management

- Process - fundamental concept in OS
 - Process is an instance of a program in execution.
 - Process needs resources - CPU time, memory, files/data and I/O devices.
- OS is responsible for the following process management activities.
 - Process creation and deletion
 - Process suspension and resumption
 - Process synchronization and interprocess communication
 - Process interactions - deadlock detection, avoidance and correction

OS Task: Memory Management

- Main Memory is an array of addressable words or bytes that is quickly accessible.
- Main Memory is volatile.
- OS is responsible for:
 - Allocate and deallocate memory to processes.
 - Managing multiple processes within memory - keep track of which parts of memory are used by which processes. Manage the sharing of memory between processes.
 - Determining which processes to load when memory becomes available

OS Task: Secondary Storage and I/O Management

- Since primary storage is expensive and volatile, secondary storage is required for backup
- Disk is the primary form of secondary storage.
 - OS performs storage allocation, free-space management and disk scheduling.
- I/O system in the OS consists of
 - Buffer caching and management
 - Device driver interface that abstracts device details
 - Drivers for specific hardware devices

OS Task: File System Management

- File is a collection of related information defined by creator - represents programs and data
- OS is responsible for
 - File creation and deletion
 - Directory creation and deletion
 - Supporting primitives for file/directory manipulation.
 - Mapping files to disks (secondary storage).
 - Backup files on archival media (tapes)

OS Task: Protection and Security

- Protection mechanisms control access of programs and processes to user and system resources.
 - Protect user from himself, user from other users, system from users.
- Protection mechanisms must:
 - Distinguish between authorized and unauthorized use.
 - Specify access controls to be imposed on use.
 - Provide mechanisms for enforcement of access control.
 - Security mechanisms provide trust in system and privacy
 - authentication, certification, encryption etc.

OS Task: Networking

- Connecting processors in a distributed system
- Distributed System is a collection of processors that do not share memory or a clock.
- Processors are connected via a communication network.
- Advantages:
 - Allows users and system to exchange information
 - provide computational speedup
 - increased reliability and availability of information.